

Workshop

NLP using Generative AI

Module 3

Speaker


Amar Lalwani

Building a startup (stealth mode)
Ex - Head of AI and R&D, funtoot




Module 3

Diving into GPT

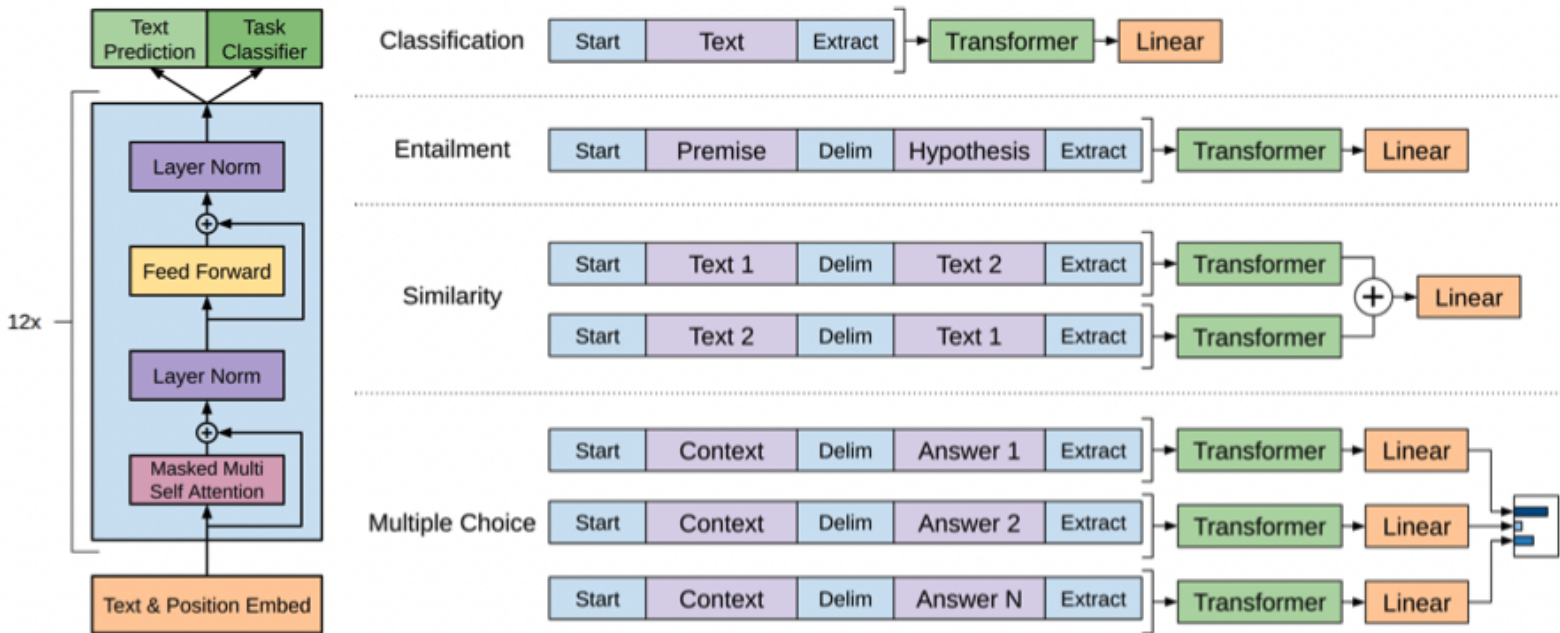


Sophisticated Next Word Predictors



Improving Language Understanding by Generative Pre-Training

- followed by (task-agnostic) Discriminative Fine-Tuning



https://s3-us-west-2.amazonaws.com/openai-assets/research-covers/language-unsupervised/language_understanding_paper.pdf

- Pre-training transformer language models
 - Task specific fine-tuning
 - Eliminating need of task specific architectures
 - Limitation: while architecture is task-agnostic, still task specific datasets and fine-tuning needed
- For each task, large high quality dataset needed. May not be possible always
- Out of distribution generalization decays with
- ◆ Expressiveness of the model
 - ◆ Narrowness of the training distribution
- Humans to do not require large training datasets to learn most natural language tasks
- ◆ Brief directive in natural language or at-most a tiny number of demonstrations is enough
 - ◆ This allows humans to seemingly mix together or switch between various tasks and skills

GPT-3: Few-shot Learner

The three settings we explore for in-context learning

Zero-shot

The model predicts the answer given only a natural language description of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 cheese => ..... ← prompt
```

One-shot

In addition to the task description, the model sees a single example of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 sea otter => loutre de mer ← example
3 cheese => ..... ← prompt
```

Few-shot

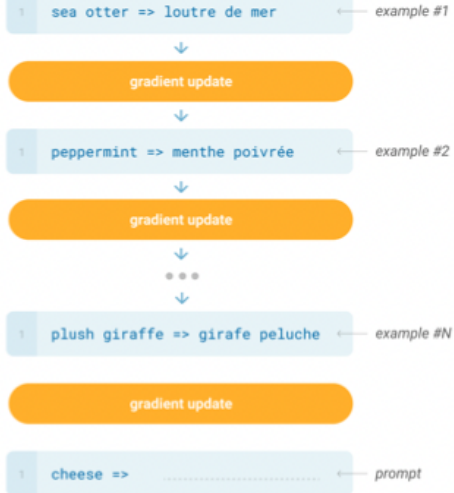
In addition to the task description, the model sees a few examples of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 sea otter => loutre de mer ← examples
3 peppermint => menthe poivrée ←
4 plush girafe => girafe peluche ←
5 cheese => ..... ← prompt
```


Traditional fine-tuning (not used for GPT-3)

Fine-tuning


The model is trained via repeated gradient updates using a large corpus of example tasks.





GPT-3 In-context learning: remove random symbols from a word



GPT-3: Few-shot Learner



Performance: Power Law



Let us do it!

